

Good, Fast, Cheap: How to Do Data Science with Missing Data

Matt Brems



Introduction



Managing Partner & Principal Data Scientist, BetaVector

Distinguished Faculty, General Assembly

Marketing + Comms Director, Statistics Without Borders

Previously:

Data Science @ Optimus Consulting

Enterprise Analytics @ Smucker's

M.S. Statistics @ The Ohio State University

Recommended Reads:

Data-Driven Thinking: "Factfulness"

Data Visualization: "Storytelling with Data"

Data Science: "An Introduction to Statistical Learning with Applications in R"

Agenda

1. **Introduction to missing data.**
2. Strategies for doing data science with missing data.
 - a. Avoid missing data.
 - b. Ignore missing data.
 - c. Account for missing data.
 - i. Unit missingness.
 - ii. Item missingness.
3. Practical considerations and warnings.

How big of a problem is missing data?

This is a difficult question to answer.


Practically, we can only see what we observe.

We can use **simulated data** to help answer this question.

To the notebook!

`00_interactive_plot.ipynb`

What is a realistic approach for us?



Source: <http://www.pyragraph.com/2013/05/good-fast-cheap-you-can-only-pick-two/>

What is a realistic approach for us?



Fast and Cheap Analysis: Drop all missing values or single imputation.

Good and Cheap Analysis: Proper imputation or pattern submodel method.

Good and Fast Analysis: Gather data in a complete manner.

Agenda

1. Introduction to missing data.
2. **Strategies for doing data science with missing data.**
 - a. Avoid missing data.
 - b. Ignore missing data.
 - c. Account for missing data.
 - i. Unit missingness.
 - ii. Item missingness.
3. Practical considerations and warnings.

Avoiding missing data.

- It's often more expensive up front but **cheaper in the long run to avoid missing data** than to make guesses about how to best handle our missing data.
- Decrease burden on your respondent.
- Change method of data collection.
- Improve accessibility.
- Change timing of your survey.
- Minimize length of questionnaire.
- Consider content of your survey.

Ignoring missing data.

- We **assume** that observations for which we've observed data are **similar to observations for which we're missing data**.
- One *general, **very rough** guideline* is that we *may* be OK ignoring missing data if less than 5% of our data is missing.
 - If we're doing supervised learning and we're missing a lot of our Y variable, this may be inadvisable.
 - If we're missing a lot from meaningful variables, this may be inadvisable.

Before “accounting for” missing data, let’s shift our mindset.

- There’s a naive belief that we can just plug in the gaps in our data.
 - This is known as **imputation**.
 - We have to do this in a specific way, or we’re just making up data.
- **In most cases, we aren’t “fixing” data. We’re just learning how to cope with it!**

Agenda

1. Introduction to missing data.
2. Strategies for doing data science with missing data.
 - a. Avoid missing data.
 - b. Ignore missing data.
 - c. **Account for missing data.**
 - i. Unit missingness.
 - ii. Item missingness.
3. Practical considerations and warnings.

What's the difference: Unit vs. Item Missingness?

Unit missingness has all values missing from an observation.

- Index 3.

Item missingness is where some, but not all, values are missing from an observation.

- Indices 1, 2, and 10,000.

Index	Age	Sex	Income
1	NA	M	NA
2	39	NA	75000
3	NA	NA	NA
4	28	F	50000
...
10000	18	F	NA

To the notebook!

`01_unit_missingness.ipynb`

Pulling the pieces together in one workflow...

1. I evaluate how much missing data I have during EDA. **Is it worth my time to try to address it?**
2. Is it reasonable to attempt **deductive imputation**?
3. If my goal is to **generate predictions**, then use the **pattern submodel** method.
4. If my goal is to **conduct inference**, then use the best imputation method available.
 - a. If I'm doing proper imputation, combine output and results using **Rubin's rules**.

Scenario 1: Missing Completely at Random (MCAR)

I'm a grad student in a lab. While pipetting, I reach for my pen but accidentally knock a Petri dish off of the desk. From this Petri dish, I lose the data that I otherwise would have collected.

This is called **missing completely at random**.

- The data of interest is not systematically different between missing and observed.

bacteria on day 1	bacteria on day 2
10mm	15mm
12mm	12mm
9mm	11mm
10mm	11mm
15mm	19mm
13mm	15mm
11mm	16mm

Scenario 2: Missing at Random (MAR)

I work for the Department of Transportation. A sensor on the Pennsylvania Turnpike broke and did not gather information between 7:00am and 10:00am.

This is called **missing at random**.

- **Conditional on data we have observed**, the data of interest is not systematically different between missing and observed.
- Whether or not a data point is missing is dependent on observed data.

time	number of vehicles
4:00	206
5:00	519
6:00	934
7:00	1,650
8:00	1,921
9:00	1,010
10:00	889

Scenario 3: Not Missing at Random (NMAR)

I administer a survey with a question about income. Those who have lower incomes are less likely to reply to the income question.

This is called **not missing at random**.

- The data of interest are systematically different for missing and observed.
- Whether or not an observation is missing depends on the value of the unobserved data itself!

id	income
A	48,000
B	35,000
C	105,000
D	62,000
E	80,000
F	50,000
G	75,000

How do we account for item nonresponse?

- Deductive Imputation
- Mean/Median/Mode
- Single Regression Imputation
- Proper Imputation
- Pattern Submodel Approach



Deductive Imputation

- We use logical rules to fill in missing values.
 - Survey asks if the respondent was the victim of a crime in the last twelve months.
 - Respondent says no.
 - Survey then asks if the respondent was the victim of a violent crime in the last twelve months.
 - Respondent leaves this answer blank.
- Requires specific coding, can be time consuming, requires no inference, can be used regardless of missingness type.

Mean/Median/Mode Imputation

- For any “NA” value in a given column, replace “NA” with the mean, median, or mode.
- Quick, easy to implement, seems reasonable, can significantly distort histogram, underestimates variance, should only be considered if data is MCAR.

Mean Imputation




Why is underestimating variance a bad thing?

$$\begin{aligned}s &= \sqrt{\frac{1}{k} \sum_{i=1}^k (x_i - \bar{x})^2} \\&= \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2} \\&= \sqrt{\frac{1}{n} \left[\sum_{i=1}^k (x_i - \bar{x})^2 + \sum_{i=k+1}^n (\bar{x} - \bar{x})^2 \right]} \\&= \sqrt{\frac{1}{n} \sum_{i=1}^k (x_i - \bar{x})^2}\end{aligned}$$

Why is underestimating variance a bad thing?

- If we underestimate variance in a confidence interval, our **confidence interval gets smaller for the same level of confidence!**
 - Our results ***artificially*** look more precise... but only because we imputed the mean!
- If we underestimate the variance in a hypothesis test, **our p -value will *artificially* get smaller.**
 - Our p -value may look significant... but only because we imputed the mean!



Mode Imputation



Single Regression Imputation

- Train a model on the rows of your data that are **fully observed**.
 - Suppose I'm missing income for some people, but have observed age and highest level of education for everybody.
 - X = age and highest level of education, Y = income.
 - Fit a model.
- For any “NA” value in a given column, replace “NA” with predicted value from that model.
- Seems reasonable, still distorts histogram, underestimates variance, should only be considered if data is MCAR or MAR.

Single Regression Imputation



Single Regression Imputation




Figure 25.2 *Deterministic and random imputations for the 241 missing values of earnings in the Social Indicators Survey. The deterministic imputations are exactly at the regression predictions and ignore predictive uncertainty. In contrast, the random imputations are more variable and better capture the range of earnings in the data. See also Figure 25.1.*

Proper Imputation


- In order to properly impute data, we need to fill in missing values with the proper amount of certainty/uncertainty.
- Replacing an NA with one value treats things like we know the true value.
- Therefore, **we need to impute multiple times.**

Proper Imputation

- We will make multiple copies (say 10) of our dataset.
- We will use random regression imputation to generate one value for each NA in each dataset.
- Once each of our 10 datasets are complete, we will do our “final model” or “final analysis” on each dataset.
- We will then combine the results of our multiple models together, just like we aggregate results in an ensemble model.

Proper Imputation

- We will make multiple copies (say 10) of our dataset.
- We will use random regression imputation to generate one value for each NA in each dataset.
- Once each of our 10 datasets are complete, we will do our “final model” or “final analysis” on each dataset.
- We will then combine the results of our multiple models together, just like we aggregate results in an ensemble model.



Proper Imputation & Rubin's Rules

- **If you're generating predictions**, you can just average your predictions together in a regression problem or select the plurality class in a classification problem.
- **If your goal is to do inference** (e.g. understand how X affects Y) and are fitting a linear model to each dataset, then you get a slope and y-intercept for each model.
 - There are a set of tools, called **Rubin's rules**, that will allow you to take the slopes and y-intercepts from each model and combine them together.
 - Check out documentation in the repo if interested!

To the notebook!

`02_item_missingness.ipynb`

How do we account for item nonresponse?

- Deductive Imputation
- Mean/Median/Mode
- Single Regression Imputation
- Proper Imputation
- Pattern Submodel Approach

Pattern Submodel Approach

- Big Picture: We will break our dataset into subsets based on missingness pattern. We will then build one model on each subset, creating many different models.

Pattern Submodel Approach

Y	X1	X2
obs	obs	obs
obs	obs	obs
obs	obs	NA
obs	obs	NA
obs	NA	obs
obs	NA	obs
obs	NA	NA

- Big Picture: We will break our dataset into subsets based on missingness pattern. We will then build one model on each subset, creating many different models.

Pattern Submodel Approach

- The pattern submodel method is expected to outperform imputation methods when data are NMAR, and will perform about as well as imputation methods when data are MCAR or MAR.
- You can generate predictions for test observations containing missing data.
- **It does not require missingness assumptions!**
- This is not a well-understood method for inference.

Pulling the pieces together in one workflow...

1. I evaluate how much missing data I have during EDA. **Is it worth my time to try to address it?**
2. Is it reasonable to attempt **deductive imputation**?
3. If my goal is to **generate predictions**, then use the **pattern submodel** method.
4. If my goal is to **conduct inference**, then use the best imputation method available.
 - a. If I'm doing proper imputation, combine output and results using **Rubin's rules**.

Agenda

1. Introduction to missing data.
2. Strategies for doing data science with missing data.
 - a. Avoid missing data.
 - b. Ignore missing data.
 - c. Account for missing data.
 - i. Unit missingness.
 - ii. Item missingness.
3. **Practical considerations and warnings.**

Warning 1

- If your goal is just to have a “complete” data set for further analysis, **be very careful!**
 - After you construct this dataset, **nobody will know the difference between observed and imputed data.**

Warning 2

- **Improper imputation is making up data.**
 - This is an **unethical** thing to do.
 - Avoid doing this!
- **Proper imputation** is not making up data.
 - Proper multiple imputation will impute data while properly estimating the amount of variance in your data.

Warning 3

- There are `sklearn` methods (`SingleImputer`, `IterativeImputer`) that are new and that can be used.
- However, `IterativeImputer` is currently experimental!
- **Proceed with caution.**

What is a realistic approach for us?



Fast and Cheap Analysis: Drop all missing values or single imputation.

Good and Cheap Analysis: Proper imputation or pattern submodel method.

Good and Fast Analysis: Gather data in a complete manner.

THANK YOU



LinkedIn: Matthew Brems

Twitter: @matthewbrems

Github: matthewbrems

Email: matt@betavector.com



hello@betavector.com